AI in the Real World

Data, Power, & Society - Day 21

Case Design: Scaffold 1 Feedback

CMSE 492 (aka CMSE 101) • Fall 2026

Prof. Danny Caballero

Today's Six Digits: 930155

CMSE 492 Attendance

Scaffold 1: Your Feedback Is In

- A **summary** is on D2L with a link to your Google Doc.
- **2–3 comments** are attached to specific passages in your Google Doc.
- **3 of 8 groups** met the standard on the first submission. *Most of the rest have one or two small gaps.*
- **Revise Part 1 in place** before you build Scaffold 2 on top of it.

Revising is a normal part of this course.

How I Wrote Your Feedback

- **I read every scaffold and took my own notes** on each one.
- I used **Claude for research into your cases**: checking that sources say what you claim, and finding laws, studies, and products relevant to your topic. *Note: I might get things wrong!*
- I used **Claude to edit my feedback** so it's clear and specific. *This is to make sure you focus on specific changes.*
- **The judgments are mine:** whether you met the standard, and what your design should focus on.

Verify before you cite. If a comment points you to a source, open it and confirm it says what's claimed. Citing it is still your responsibility.

What We Looked For

A scaffold meets the standard when it has:

- **≥750 words** in 1.1–1.4, written in paragraphs
- **≥2 cited sources with specific numbers** in 1.2
- At least one **low-power stakeholder**, plus a **priority paragraph**
- **DTPA paragraphs specific to your case**, not to AI in general
- A **research map** of 4–6 *distinct* sources of ≥2 kinds
- **Every section present**, including References and Team Roles
- **All Roles assigned** and including all group members

What Went Well

Some moves we want to see more of:

- **Critique vs. constraint:** *"Flock cut its default retention … but that is a company setting, not a law."*
- **Turning a rate into people:** a 1% false-positive rate × 75,000 papers ≈ **750 students**
- **Naming who's invisible:** *"Civilians are largely invisible in this data."*
- **Separating harms:** one rule fixes bad alerts but not a searchable location history
- **Drawing a boundary:** the device should *not* tell someone it's safe to cross the street

Pattern 1: Sources Didn't Always Hold Up

- **1.1 was uncited almost everywhere**, even though your biggest claims seem to appear there.
- **Research map and References sometimes didn't match** in 6 of 8 scaffolds. *Check your references.*
- **Check that the source says it.** We found:
  - findings credited to papers that don't contain them
  - a reference crediting an author who doesn't exist (it was the *system's* name)
  - incidents retold with the wrong details
- **Who is speaking?** News about a study isn't research. A company selling an AI "humanizer" has a stake in saying detectors fail.

Pattern 2: Numbers Stretched in the Retelling

Small wording changes can change your argument:

- "51.6% **of patients**" → it was 51.6% **of emergency cases**
- "AI as an **alternative**" → the survey asked about AI as an **additional** resource
- "Palantir reported **80,000** users" → **20,000** reported; 80,000 was a researcher's **estimate**

> **For every number: what was measured, on what, and compared to what?**
>
> Sometimes, I think this was because someone didn't review the entire document.

Pattern 3: Practices and Actions Drifted

- Some Practices and Actions paragraphs described **LLMs or AI in general**, not your case.
- One Practices paragraph **restated the Tools paragraph** in different words.
- **Practices** is about *people and institutions*: who builds, buys, operates, and oversees it, whose labor it depends on, and which rules apply.
- **Actions** is about *what it decides*: who benefits, who pays, and what it rewards.

Same lesson as Case Study I: these are likely the hardest to research, and yet the most important to get specific.

Small Things That Matter

- Delete leftover **template prompts** and **placeholders** ("[Working Title]")
- **Complete Team Roles:** every lead named, plus a **coordinator**
- **Update your working title** to your actual case or proposal
- **Read the whole document together** once before you submit

Scaffold 2: Find Your Most Important Lens(es)

One Lens Drives Your Design

Your feedback names **one DTPA component** (sometimes two) as the most central to your case.

- **All four sections** must still meet the standard.
- That lens gets your **deepest analysis** (about 400 words instead of 300).
- Your **design direction (2.5)** should grow out of it.

How we picked the lens(es): **where does the harm or the value actually come from?**

Which Lens?

- **Data:** the failure was stale, missing, or unconsented data
  - *e.g., database records that didn't change; improperly collected or reused data; data that doesn't match the effort*
- **Tools:** the value depends on error rates
  - *e.g., is it better than using a cane, or than guessing "pass" every time?*
- **Practices:** the harm is misuse of access, or no one is accountable
  - *e.g., searches no one authorized; operating outside the legal frameworks that are playing catch up*
- **Actions:** the harm is a decision the system makes
  - *e.g., how urgent is this? Is this score proof of cheating? could a well-trained person do this better?*

Pathway I: What's Already Out There?

- **Survey what exists:** products, apps, and how people do this work today.
- Say what your design does **differently**: cost, focus, privacy, who it serves.
- **Compare to a baseline:** a simple guess, human performance, or the current approach.
- Say what your tool **deliberately does not do**.

"Existing product, plus more features" isn't a design direction.

Pathway II: What Does Your Restriction Add?

Most Pathway II proposals **overlap with rules that already exist**:

- state bills and federal laws
- agency directives and procedures
- the company's **own disclaimers and policies**

Scaffold 2 must say:

1. What **exactly** your restriction requires or forbids
2. What it **adds** beyond existing rules
3. **Who enforces it**, and what happens when it's violated

> If the company already says it, it isn't a restriction yet.

What's Different in the Scaffold 2 Template

- **Response to feedback:** choose your **three most important changes**. For each: the feedback, what you changed and why, and where it appears.
- **No separate "revise Part 1" section:** revise Part 1 **in place**.
- **Annotated bibliography = your reference list.** Label each source's *kind*. Every in-text citation appears there, and vice versa.
- **Error figures** must be compared to a **baseline**. *Here, similar or historical cases might help.*

Today

1. **Read** your Scaffold 1 summary and comments together *(~10 min)*
2. **Discuss:** do you agree with your DTPA focus? *Disagreeing is fine; say why in your response.*
3. **Plan:** who revises Part 1, and who leads your focus lens
4. **Then:** keep working on **Full Case II (Part B)**

We'll be circulating. Ask us anything.

Full Case II Update: Graded by Wednesday

- Feedback on your **Full Case II Update (Part A)** will be in your Google Doc **by Wednesday's class (Oct 21)**.
- Use it to finish **Part B**, due **Tues, Oct 27**.

No Exit Ticket Today

- Exit Tickets are **only after Wednesday's class**.
- **No Exit Ticket today.**

Reminders

- **Wed, Oct 21:** Revised Case Study I due (if needed)
- **Wed, Oct 21:** Full Case II Update feedback in your Doc by class
- **Thu, Oct 22:** Week 8 Forum Replies (optional)
- **Mon, Oct 26 – Tue, Oct 27:** Fall Break, no class
- **Tue, Oct 27:** Full Case II Complete (Part B) due
- **Tue, Oct 27:** Week 8 Individual Effort & Metacognitive Report due
- **Tue, Oct 27:** Week 9 Forum Post due
- **Fri, Nov 6:** Case Design Scaffold 2 due (with revised Part 1 and response to feedback)